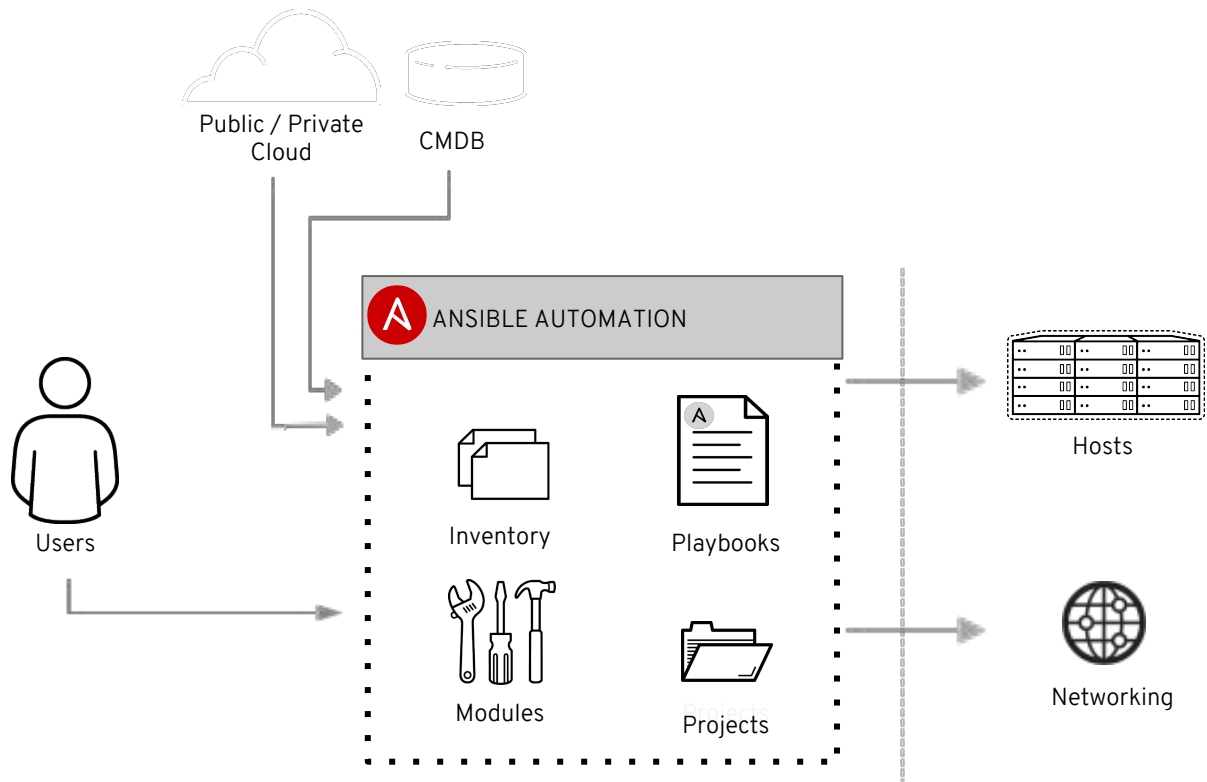


ANSIBLE

ANSIBLE BEST PRACTICES: THE ESSENTIALS

Magnus Glantz, Senior Solution Architect
Red Hat

GitHub/freenode: mglantz






THE ANSIBLE WAY

Principal 1 - COMPLEXITY KILLS PRODUCTIVITY

ANSIBLE

Now

We need to automate



No problem. I'm Phil. I'll have this automated in a bit.



5 years later



Phil wrote that, good luck, he got killed in the great dog-regexp war of 2019.



Heey... Do you know what

```
^(?:((?:0?[13578]|1[02])(M-|L)31)1|((?:0?[13-9]|1[0-2])(M-|L)(?:29|30|2))((?:1[6-9]|2-9)d)?\d{2})$|^((?:0?2(M-|L)29|3((?:1[6-9]|2-9)d)?((?:0[48]|[2468]|[048]|[13579][26])|((?:16|[2468][048][3579][26])00))))$|^((?:0?[1-9])|((?:1[0-2])(M-|L)(?:0?[1-9]|1\d|2[0-8])\4((?:1[6-9]|2-9)d)?\d{2}))$
```

...does?

Principal 2 - OPTIMIZE FOR READABILITY

ANSIBLE

```
1  - name: Install Tomcat Application server and deploy sample Java app
2    hosts: all
3    tasks:
4      - name: Ensure tomcat is installed
5        yum:
6          name: tomcat
7          state: present
8      - name: Ensure tomcat service is enabled and started
9        service:
10         name: tomcat
11         enabled: yes
12         state: started
13     - name: Download and deploy Java application
14       get_url:
15         url: https://tomcat.apache.org/tomcat-6.0-doc/appdev/sample/sample.war
16         dest: /var/lib/tomcat/webapps/sample.war
17         mode: 0777
```



Principal 3 - THINK DECLARATIVELY

ANSIBLE

Ansible is a desired state engine by design. If you're trying to "write code" in your plays and roles, you're setting yourself up for failure. Our YAML-based playbooks were never meant to be for programming.


Treat your Ansible content like code

- Version control your Ansible content
- Start as simple as possible and iterate
 - Start with a basic playbook and static inventory
 - Refactor and modularize later



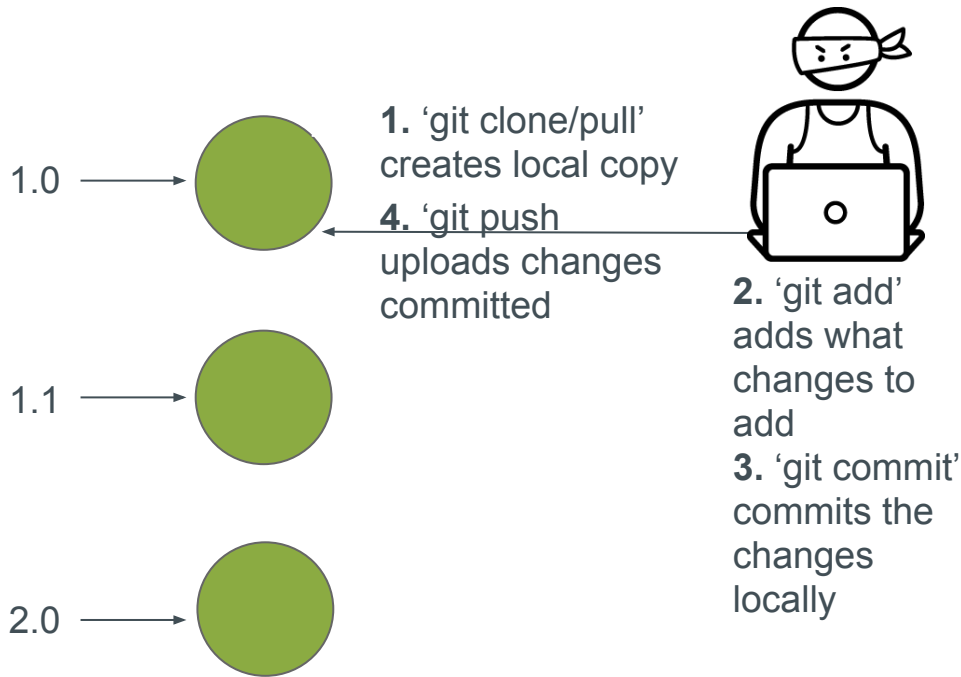
Treat your Ansible content like code

1. **Ansible** doesn't require version control
2. **When** you scale out your Ansible usage (aka. automate all things) you'll have many different teams collaborating
3. **Version control** was invented to solve common collaboration challenges
4. **Git** has earned its worldwide popularity the hard way and is in the core of many of the world's most popular collaboration services and products




Example: Version control

1. **A** git repository stores files
2. **Access** controls are specific to repositories
3. **All** changes to all files are tracked
4. **When** you want to make a change to a file you first make a local copy of the repository which is stored on your computer, you then change the file locally, commit the change locally and then go ahead and tell git to copy this local change to the repository.




Example: GitHub workflow

1. **Does not** require GitHub, the workflow model is just called that
2. **A** very simple workflow
3. **Master** branch is always possible to release
4. **Branches** are where you develop and test new features and bugfixes.
5. **Yes**, I wrote test. If you do not test your Ansible code you cannot keep the master branch releasable and this all fails.



Treat your Ansible content like code



Do It with Style

- Create a style guide for developers
- Consistency in:
 - Tagging
 - Whitespace
 - Naming of Tasks, Plays, Variables, and Roles
 - Directory Layouts
- Enforce the style
- Check out ansible-lint

I deserve
another demo.



Implement a test framework for playbooks

A basic framework for Ansible testing is:

- Verify correct syntax with
 - a. `ansible-playbook --syntax-check your-playbook.yml`
- Verify style for bad practices and behaviour that could potentially be improved
 - a. `ansible-lint your-playbook.yml`
- Run your playbook or role and ensure it completes without failures.
- Run your playbook or role again and ensure that no changes are reported, this ensures playbook idempotency, a key feature of Ansible.
- Query your application's API or do another external test of it's functionality.
- Implement your testing framework into a CI/CD pipeline for your playbooks

Read more

<https://github.com/mglantz/ansible-roadshow/tree/master/labs/lab-9>

```
basic-project
├── inventory
│   ├── group_vars
│   │   └── web.yml
│   ├── host_vars
│   │   └── db1.yml
│   └── hosts
└── site.yml
```

```
myapp
├── roles
│   ├── myapp
│   │   ├── tasks
│   │   │   └── main.yml
│   │   └── ...
│   ├── nginx
│   │   └── ...
│   └── proxy
│       └── ...
└── site.yml
```

myapp

```
|— config.yml
|— provision.yml
|— roles
|   └─ requirements.yml
└─ site.yml
```


Give inventory nodes human-meaningful

EXHIBIT A

```
10.1.2.75  
10.1.5.45  
10.1.4.5  
10.1.0.40
```

```
w14301.example.com  
w17802.example.com  
w19203.example.com  
w19304.example.com
```



EXHIBIT B

```
db1  ansible_host=10.1.2.75  
db2  ansible_host=10.1.5.45  
db3  ansible_host=10.1.4.5  
db4  ansible_host=10.1.0.40
```

```
web1  ansible_host=w14301.example.com  
web2  ansible_host=w17802.example.com  
web3  ansible_host=w19203.example.com  
web4  ansible_host=w19203.example.com
```

Group hosts for easier inventory selection and less conditional tasks -- the more groups the better.

WHAT

```
[db]  
db[1:4]
```

```
[web]  
web[1:4]
```

```
db1 = db, east, dev
```

WHERE

```
[east]  
db1  
web1  
db3  
web3
```

```
[west]  
db2  
web2  
db4  
web4
```

WHEN



```
[dev]  
db1  
web1
```

```
[test]  
db3  
web3
```

```
[prod]  
db2  
web2  
db4  
web4
```

Use a single source of truth if you have it -- even if you have multiple sources, Ansible can unify them.

- Stay in sync automatically
- Reduce human error



The world is flat - Proper variable naming can make plays more readable and avoid variable name conflicts

- Use descriptive, unique human-meaningful variable names
- Prefix role variables with its “owner” such as a role name or package

```
apache_max_keepalive: 25
apache_port: 80
tomcat_port: 8080
```

```
- name: Clone student lesson app for a user
host: nodes
tasks:
  - name: Create ssh dir
    file:
      state: directory
      path: /home/{{ username }}/.ssh

  - name: Set Deployment Key
    copy:
      src: files/deploy_key
      dest: /home/{{ username }}/.ssh/id_rsa

  - name: Clone repo
    git:
      accept_hostkey: yes
      clone: yes
      dest: /home/{{ username }}/exampleapp
      key_file: /home/{{ username }}/.ssh/id_rsa
      repo: git@github.com:example/apprepo.git
```

EXHIBIT A

- Embedded parameter values and repetitive home directory value pattern in multiple places
- Works but could be more clearer and setup to be more flexible and maintainable

```
- name: Clone student lesson app for a user
  host: nodes
  vars:
    user_home_dir: /home/{{ username }}
    user_ssh_dir: "{{ user_home_dir }}/.ssh"
    deploy_key: "{{ user_ssh_dir }}/id_rsa"
    app_dir: "{{ user_home_dir }}/exampleapp"
  tasks:
    - name: Create ssh dir
      file:
        state: directory
        path: "{{ user_ssh_dir }}"

    - name: Set Deployment Key
      copy:
        src: files/deploy_key
        dest: "{{ deploy_key }}"

    - name: Clone repo
      git:
        dest: "{{ app_dir }}"
        key_file: "{{ deploy_key }}"
        repo: git@github.com:example/exampleapp.git
        accept_hostkey: yes
        clone: yes
```

EXHIBIT B

- Parameters values are set thru values away from the task and can be overridden.
- Human meaningful variables “document” what’s getting plugged into a task parameter
- More easily refactored into a role

Use native YAML syntax to maximize the readability of your plays

- Vertical reading is easier
- Supports complex parameter values
- Works better with editor syntax highlighting in editors

NO!

- `name: install telegraf`
 `yum: name=telegraf-{{ telegraf_version }} state=present update_cache=yes disabl`
 `notify: restart telegraf`
- `name: configure telegraf`
 `template: src=telegraf.conf.j2 dest=/etc/telegraf/telegraf.conf`
- `name: start telegraf`
 `service: name=telegraf state=started enabled=yes`

Better, but no

- `name: install telegraf`
`yum: >`
 - `name=telegraf-{{ telegraf_version }}`
 - `state=present`
 - `update_cache=yes`
 - `disable_gpg_check=yes`
 - `enablerepo=telegraf``notify: restart telegraf`
- `name: configure telegraf`
`template: src=telegraf.conf.j2 dest=/etc/telegraf/telegraf.conf`
- `name: start telegraf`
`service: name=telegraf state=started enabled=yes`

Yes!

```
- name: install telegraf
  yum:
    name: telegraf-{{ telegraf_version }}
    state: present
    update_cache: yes
    disable_gpg_check: yes
    enablerepo: telegraf
  notify: restart telegraf

- name: configure telegraf
  template:
    src: telegraf.conf.j2
    dest: /etc/telegraf/telegraf.conf
  notify: restart telegraf

- name: start telegraf
  service:
    name: telegraf
    state: started
    enabled: yes
```

Names improve readability and user feedback

- Give all your playbooks, tasks and blocks brief, reasonably unique and human-meaningful names

EXHIBIT A

```
- hosts: web
  tasks:
    - yum:
        name: httpd
        state: latest

    - service:
        name: httpd
        state: started
        enabled: yes
```

```
PLAY [web]
*****

TASK [setup]
*****
ok: [web1]

TASK [yum]
*****
ok: [web1]

TASK [service]
*****
ok: [web1]
```

EXHIBIT B

```
- hosts: web
  name: install and start apache
  tasks:
    - name: install apache packages
      yum:
        name: httpd
        state: latest

    - name: start apache service
      service:
        name: httpd
        state: started
        enabled: yes
```

```
PLAY [install and start apache]
*****
```

```
TASK [setup]
*****
ok: [web1]
```

```
TASK [install apache packages]
*****
ok: [web1]
```

```
TASK [start apache service]
*****
ok: [web1]
```

Focus avoids complexity

- Keep plays and playbooks focused. Multiple simple ones are better than having a huge single playbook full of conditionals
- Follow Linux principle of do one thing, and one thing well

Clean up your debugging tasks

- Make them optional with the verbosity parameter so they're only displayed when they are wanted.

```
- debug:  
  msg: "This always displays"  
  
- debug:  
  msg: "This only displays with ansible-playbook -vv+"  
  verbosity: 2
```

Don't just start services -- use smoke tests

```
- name: check for proper response
  uri:
    url: http://localhost/myapp
    return_content: yes
  register: result
  until: '"Hello World" in result.content'
  retries: 10
  delay: 1
```


Use command modules sparingly

- Use the run `command` modules like `shell` and `command` as a last resort
- The `command` module is generally safer
- The `shell` module should only be used for I/O redirect

Always seek out a module first

NO!

- name: add user
command: useradd appuser
- name: install apache
command: yum -y install httpd
- name: start apache
shell: |
systemctl start httpd && systemctl enable httpd

Yes :-)

- name: add user
user:
name: appuser
state: present
- name: install apache
yum:
name: httpd
state: latest
- name: start apache
service:
name: httpd
state: started
enabled: yes

Still using command modules a lot?

```
- hosts: all
vars:
  cert_store: /etc/mycerts
  cert_name: my cert
tasks:
- name: check cert
  shell: certify --list --name={{ cert_name }} --cert_store={{ cert_store }} | grep "{{ cert_name }}"
  register: output

- name: create cert
  command: certify --create --user=chris --name={{ cert_name }} --cert_store={{ cert_store }}
  when: output.stdout.find(cert_name) != -1
  register: output

- name: sign cert
  command: certify --sign --name={{ cert_name }} --cert_store={{ cert_store }}
  when: output.stdout.find("created") != -1
```

Develop your own module

```
- hosts: all
  vars:
    cert_store: /etc/mycerts
    cert_name: my cert
  tasks:
    - name: create and sign cert
      certify:
        state: present
        sign: yes
        user: chris
        name: "{{ cert_name }}"
        cert_store: "{{ cert_store }}"
```

- Understandable by non-technical people
- CRUD (Create, read, update and delete)

Separate provisioning from deployment and configuration tasks

```
acme_corp/  
├── configure.yml  
├── provision.yml  
└── site.yml
```

```
$ cat site.yml  
---  
- import_playbook: provision.yml  
- import_playbook: configure.yml
```

Jinja2 is powerful but you needn't use all of it

- Templates should be simple:
 - Variable substitution
 - Conditionals
 - Simple control structures/iterations
 - Design your templates for your use case, not the world's
- Things to avoid:
 - Anything that can be done directly in Ansible
 - Managing variables in a template
 - Extensive and intricate conditionals
 - Conditional logic based on embedded hostnames
 - Complex nested iterations

What did we say
about
complexity?



Careful when mixing manual and automated configuration
(Or even different automation frameworks...)

- Label template output files as being generated by Ansible

```
{{ ansible_managed | comment }}
```

Keep in mind

- Like playbooks -- keep roles purpose and function focused
- Use a `roles/` subdirectory for roles developed for organizational clarity in a single project
- Follow the Ansible Galaxy pattern for roles that are to be shared beyond a single project
- Limit role dependencies

Tricks and tips

- Use `ansible-galaxy init` to start your roles...
- ...then remove unneeded directories and stub files
- Use `ansible-galaxy` to install your roles -- even private ones
- Use a roles files (i.e. `requirements.yml`) to manifest any external roles your project is using
- Always peg a role to a specific version such as a tag or commit

Command line tools have their limitations

- Coordination across a distributed teams & organization...
- Controlling access to credentials...
- Track, audit and report automation and management activity...
- Provide self-service or delegation...
- Integrate automation with enterprise systems...



Complexity kills productivity
Optimize for readability
Think declaratively

Thank you

Complexity